

Week 1 part 2: Stepping; Haskell intro

Jana Dunfield

September 6, 2022

References to “Thompson” are to *Haskell: The Craft of Functional Programming*, 3rd ed.

1 Functional programming in industry

- <http://cufp.org/2015/fighting-spam-with-haskell-at-facebook.html> (Haskell at Facebook)
- <http://cufp.org/2016/yaron-minsky-keynote.html> (OCaml at Jane Street, a financial services firm)
- <http://cufp.org/2017/using-functional-programming-to-accelerate-translational-research-at-pfizer.html> (functional programming for pharmaceuticals)

An (arguably) functional programming language, Rust, is used in substantial parts of Firefox—for both reliability and performance reasons (functional programs are often easier to parallelize).

<https://tutorcs.com>

WeChat: cstutorcs

2 Stepping

The behaviour of Haskell *expressions*, such as $2 + 2$ and $(\backslash x \rightarrow x + 1) 3$, can be predicted by a process of stepping. Thompson calls stepping “calculation”; another name is “tracing”.

For very simple examples, stepping is hard to distinguish from equality in algebra:

$$\begin{array}{lcl} 2 + 2 & & 2 + 2 \\ = 4 & \Rightarrow & 4 \quad \text{by arithmetic} \end{array}$$

(Thompson uses a squiggly arrow \rightsquigarrow instead of \Rightarrow .) If you are taking notes on a laptop, I suggest typing `==>` to avoid confusion with `=>`, which is part of Haskell’s type syntax.

For slightly larger examples, however, stepping distinguishes itself from equality—we do only *one operation per step*:

$$\begin{array}{lcl} (2 + 2) + (3 + 4) & & (2 + 2) + (3 + 4) \\ = 11 & \Rightarrow & 4 + (3 + 4) \quad \text{by arithmetic} \\ & \Rightarrow & 4 + 7 \quad \text{by arithmetic} \\ & \Rightarrow & 11 \quad \text{by arithmetic} \end{array}$$

And whereas mathematical equality may have no particular goal—we can write $5 = 1 + (6 - 2) = 3 + 2 = 2 + 1 + 2$ —stepping moves towards a result. The Haskell expression `11` will not step “backwards” to `4 + 7`.

Adding integers may seem simple, but it is not always clear how stepping should work. How should `1 / 3` step? In a previous iteration of 360, I received four suggestions:

1. $(1 / 3)$
 $\Rightarrow 0$ by arithmetic
2. $(1 / 3)$
 $\Rightarrow 0.3333333333333333$ by arithmetic
3. $(1 / 3)$
 $\Rightarrow 1$ by arithmetic
4. $(1 / 3)$
 $\Rightarrow \frac{1}{3}$ by arithmetic

All of these are reasonable choices, and at least three of them are found in actual programming languages; the operator `/` exists in most languages, but they don’t agree on its meaning (its *semantics*).

- Suggestion 1 interprets `/` as integer division, rounding down. C, Java and Python interpret / this way.

(In algorithms textbooks, this would be written $\lfloor 1 / 3 \rfloor$: mathematicians usually assume / means division on real numbers, so they have to “floor” or “truncate” the result $\frac{1}{3}$ by applying the operation $\lfloor \dots \rfloor$.)

- Suggestion 2 interprets `/` as real (floating-point) division. Haskell interprets / this way.

- Suggestion 3 interprets `/` as integer division, rounding *up* (like $\lceil 1/3 \rceil$). I don't know of any language that does this, but there probably are some.
- Suggestion 4 interprets `/` as exact division or rational division ($1/3$ being the ratio of two integers). The Racket language interprets `/` this way (though in Racket, you write `(/ 1 3)` rather than `1 / 3`). A student said that Matlab does this, and I expect that Maple and Mathematica do as well.

Aside:

```
/* example showing the semantics of the / operator in Java */
public class hmm {
    public static void main (String[] args) {
        System.out.println (1 / 3);          /* prints 0 */

        System.out.println (3 * (1.0 / 3));  /* prints 1.0 */
    }
}
```

3 Using GHCi's read-eval-print loop

We interact with GHCi through its “read-eval-print loop” (which some people call a REPL, pronounced “repple”). Read-eval-print loops are available for many languages, including Python, but not others (including Java, C and C++).

The read-eval-print loop is:

1. **read** input
2. **evaluate** input
3. **print** result
4. loop back to 1

When you enter an expression such as `2 + 2` into GHCi, the loop **reads** the expression, **evaluates** it to a result (namely 4), **prints** the result, and then waits to read more input.

(If you are familiar with command-line shells such as `bash`, they are generally not “read-eval-print”, but “read-run”: by default, they don't print the exit status. If you are not familiar with command-line shells, don't worry about this.)

Typing large expressions is tedious and unreliable, however, so most of the time we will enter Haskell code in files (with the extension `.hs`), load them into GHCi (as described below), and then evaluate some expressions. For example, if we define a function `triple` in a `.hs` file, we can type

```
triple 10
```

and (hopefully) get the result 30.

3.1 Loading files

To load `filename.hs` into GHCi, enter a *load command*:

```
:load filename.hs
or
:load filename
or
:l filename
```

Once the file is loaded, the prompt “Prelude>” becomes “*Main>”. (If the file contains a module, then the module’s name will be printed instead of “Main”.)

A load command is *not* a Haskell expression: When you enter a line beginning with `:` into GHCi, it interprets the line as a command. Commands ask GHCi to take some action that is not “evaluate a Haskell expression”. You can see a list of GHCi commands by entering `:?` or `:help`, though I find this more useful for reminding myself of commands’ spelling than for learning how they work.

The `:reload` command, which can be abbreviated `:r`, reloads the previously loaded file. So if you change `filename.hs` and want to use the updated version, enter `:r`.

3.2 Seeing the type of an expression

Like Java, but unlike Python, Haskell is a “statically typed” language. You can ask GHCi for the type of an expression by entering the command

```
:type Expression
```

Haskell’s type system prioritizes power (being good at catching subtle bugs) over ease of learning, so this can give surprising results:

```
*Main> :type 3
3 :: Num a => a
```

I won’t explain this in full detail now, but roughly, this is Haskell “hedging its bets”: the expression `3` can have a variety of numeric types (including integers and floating-point numbers, which behave differently), so Haskell refuses to choose just one. Instead, it says `3` has type ‘`a`’, where all we know is that `a` is some numeric type (‘`Num a`’).

If we tell Haskell that we want a specific numeric type, say `Int`, it accepts our request and says the expression `(3 :: Int)` has type `Int`.

```
*Main> :type (3 :: Int)
(3 :: Int) :: Int
```

If `3` is an `Int`, then Haskell decides (reasonably) that adding `4` to it gives an `Int`, even though we wrote `4` and not `(4 :: Int)`:

```
*Main> :type (3 :: Int) + 4
(3 :: Int) + 4 :: Int
```

This is all somewhat confusing, but Haskell is more decisive when you ask it about definitions in loaded files: if you put

```
x = 3
```

in a file `a.hs`, load `a.hs`, and then ask for the type of `x`, you get a clearer answer. Try it!

4 Functions

An anonymous function, called a *lambda* (λ), is written in Haskell by writing a backslash (which sort of resembles λ), followed by the name of the argument (called the *bound variable*), followed by \rightarrow , followed by the function body:

```
\x -> x + 1
^      ^^^^^
bound  body
variable
```

We can call this function by writing it, then giving an argument:

```
(\x -> x + 1) 6
```

This is a small expression: try it in the GHCi read-eval-print loop.

How does this step? Since GHCi prints 7 at the end, there must be a sequence of steps

Assignment Project Exam Help
<https://tutorcs.com>
It seems plausible that the last step is to add 6 and 1:

WeChat: cstutorcs
$$\begin{aligned} & (\lambda x \rightarrow x + 1) 6 \\ \Rightarrow & \vdots \\ \Rightarrow & 6 + 1 \\ \Rightarrow & 7 \end{aligned} \quad \text{by arithmetic}$$

In fact, there are just two steps: the second is arithmetic, the first is a function application (another name for a function call).

$$\begin{aligned} & (\lambda x \rightarrow x + 1) 6 \\ \Rightarrow & 6 + 1 && \text{by function application} \\ \Rightarrow & 7 && \text{by arithmetic} \end{aligned}$$

If we can step from $(\lambda x \rightarrow x + 1) 6$ to $6 + 1$ there must be a stepping rule that explains how to do this.

4.1 Stepping rule for function application

If f is a function,
and the bound variable of f is x ,
and the body of f is $body$, then

$$f \text{ arg} \Rightarrow \text{body with arg substituted for } x$$

Read the last line as:

f applied to argument arg steps to $body$ with arg substituted for x

■ **Exercise 1.** Using the stepping rule for function application, then the stepping rule for arithmetic, fill in the following:

$$\begin{aligned} & (\backslash z \rightarrow 8 - z) 0 \\ \Rightarrow & \hspace{10em} \text{by function application} \\ \Rightarrow & \hspace{10em} \text{by arithmetic} \end{aligned}$$

■ **Exercise 2.** Fill in the following:

$$\begin{aligned} & (\backslash f \rightarrow f - f) (2 + 2) \\ \Rightarrow & \hspace{10em} \text{by arithmetic} \\ \Rightarrow & \hspace{10em} \\ \Rightarrow & \hspace{10em} \end{aligned}$$

(There is only one correct *result*, zero, but several reasonable sequences of steps. Only one sequence of steps corresponds to what Haskell actually does. But I haven't yet explained what Haskell actually does!)

5 Optional readings for this lecture

- Thompson, chapter 1
 - Section 1.10, “Calculation and evaluation”, is another presentation of *stepping*.

6 Other remarks

■ **Remark 1.** In Thompson's discussion of fixed-size integers, the bound 2147483647—which is $2^{31} - 1$ —is almost certainly out of date. To see the value on your system, enter (in GHCi)

```
maxBound :: Int
```

If you enter `maxBound` without the “`:: Int`”, you won't get anything useful. (GHCi won't even give you an error message, though Hugs will.) GHC has different maximum values for different types; when you enter `maxBound` by itself, it doesn't know what type you want.

(It's usually better to use `Integer` instead of `Int` anyway, because `Integer` avoids the possibility of overflow.)